


PAR Laboratory Assignment

Lab 4: Parallel Task Decomposition

Implementation and Analysis



UNIVERSITAT POLITÈCNICA DE CATALUNYA
FACULTAD D'INFORMÀTICA DE BARCELONA

Nuria Lou Enseñat Balaguer (nuria.lou.ensenat)
Albert Gómez Triunfante (albert.gomez.triunfante)



Index

Index	1
Introduction	2
Iterative task decomposition	3
Tile	3
Code	3
Modelfactor Analysis	4
Paraver Analysis	5
Strong Scalability	5
Finer	7
Code	7
Modelfactor Analysis	8
Paraver Analysis	9
Strong Scalability	9
Recursive task decomposition	11
Leaf	11
Code	11
Modelfactor Analysis	11
Paraver Analysis	12
Strong Scalability	13
Tree	14
Code	14
Modelfactor Analysis	14
Paraver Analysis	15
Strong Scalability	15
Summary of the elapsed execution time	17

Introduction

In this laboratory assignment, we aim to evaluate the strong scalability of different parallelisation strategies for computing the Mandelbrot set, as it presents an excellent case study for analysing parallel processing techniques due to its intricate fractal patterns generated through iterative computations in the complex plane.

The report focuses on comparing several strategies, both iterative and recursive, to understand their impact on performance and scalability. On the iterative side, we employ the tile and finer grain strategies, each offering distinct approaches to task decomposition, from coarser workloads to more granular subdivisions. For recursive methods, the leaf and tree strategies are analysed, which explore varying depths and hierarchies of task creation. These strategies provide insights into the trade-offs between load balancing, synchronisation overhead, and task management efficiency.

Our primary objective is to examine how these strategies scale with increasing thread counts while keeping the problem size constant. The analysis aims to determine which strategy—iterative or recursive—provides the best performance for this problem and why.

By exploring these different task decomposition approaches, this report seeks to provide a comprehensive understanding of the strengths and limitations of each parallelisation strategy.

Iterative task decomposition

In this section we will analyse different iterative task decompositions.

Tile

Code

For the tile version we were given a sequential approach, in order to parallelise it we had to create the necessary tasks and delete the existing variables' dependencies as we studied in the previous laboratory.

We first added the `#pragma omp task firstprivate(x, y)` as well as the `#pragma omp parallel` and `#pragma omp single` directives in order to create the necessary tasks for the tile version. Then, as we previously studied that the dependencies were caused by both variables `histogram` and `&X11_COLOR_FAKE`, in order to fix this we added a `#pragma omp atomic` and `#pragma omp critical` directives.

```
void mandel_tiled(int M[ROWS][COLS], double CminR, double CminI, double CmaxR, double CmaxI, double scale_real, double scale_imag, int maxiter)
{
    int equal;

    for (int y = 0; y < ROWS; y += TILE)
    {
        for (int x = 0; x < COLS; x += TILE)
        {
            // Set all matrix positions with the same value
            #pragma omp task firstprivate(x,y)
            {
                equal = 1;
                for (int px = x; px < x + TILE; px++) {
                    M[y][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y, scale_real, scale_imag, maxiter);
                    M[y + TILE - 1][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y + TILE - 1, scale_real, scale_imag, maxiter);
                    equal = equal & (M[y][x] == M[y][px]);
                    equal = equal & (M[y][x] == M[y + TILE - 1][px]);
                }
                for (int py = y; py < y + TILE; py++) {
                    M[py][x] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x, py, scale_real, scale_imag, maxiter);
                    M[py][x + TILE - 1] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x + TILE - 1, py, scale_real, scale_imag, maxiter);
                    equal = equal & (M[y][x] == M[py][x]);
                    equal = equal & (M[y][x] == M[py][x + TILE - 1]);
                }
                if (equal) {
                    long color = (long)((M[y][x] - 1) * scale_color) + min_color;
                    if (output2histogram) {
                        #pragma omp atomic
                        histogram[M[y][x] - 1] += (TILE*TILE);
                    }
                    for (int py = y; py < y + TILE; py++)
                        for (int px = x; px < x + TILE; px++) {
                            M[py][px] = M[y][x];
                            if (output2display) {
                                /* Scale color and display point */
                                if (setup_return == EXIT_SUCCESS) {
                                    #pragma omp critical
                                    {
                                        XSetForeground(display, gc, color);
                                        XDrawPoint(display, win, gc, px, py);
                                    }
                                }
                            }
                        }
                }
            }
        }
    }
    // Compute
    for (int py = y; py < y + TILE; py++)
        for (int px = x; px < x + TILE; px++) {
            M[py][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, py, scale_real, scale_imag, maxiter);
            if (output2histogram) {
                #pragma omp atomic
                histogram[M[py][px] - 1]++;
            }
            if (output2display) {
                /* Scale color and display point */
                long color = (long)((M[py][px] - 1) * scale_color) + min_color;
                if (setup_return == EXIT_SUCCESS) {
                    #pragma omp critical
                    {
                        XSetForeground(display, gc, color);
                        XDrawPoint(display, win, gc, px, py);
                    }
                }
            }
        }
}
```

```
main():
```

```

    #pragma omp parallel
    #pragma omp single
    mandel_tiled((int (*)(width))Hmatrix, real_min, imag_min, real_max, imag_max, scale_real, scale_imag, maxiter);

```

Modelfactor Analysis

These are the results from running the modelfactor analysis:

In this first table we can notice that when increasing the number of processors, the elapsed time is decreased until reaching its peak at 8 processors, where it does not matter adding more processors as the elapsed time will remain quite the same. We can notice the same behaviour in the SpeedUp row. But efficiency drops as we keep increasing the number of processors.

Overview of whole program execution metrics									
Number of Processors	1	2	4	6	8	10	12	14	16
Elapsed time (sec)	3.09	1.74	0.92	0.75	0.70	0.72	0.71	0.71	0.71
Speedup	1.00	1.78	3.37	4.13	4.41	4.29	4.35	4.35	4.34
Efficiency	1.00	0.89	0.84	0.69	0.55	0.43	0.36	0.31	0.27

Table 1: Analysis done on Thu Nov 21 12:15:25 PM CET 2024, par1307

In this second table we can notice that as analysed previously, the efficiency drops until 27.12% when having 16 processors. Also, the load balancing keeps decreasing due to having coarse grained tasks.

Overview of the Efficiency metrics in parallel fraction, $\phi=99.99\%$									
Number of Processors	1	2	4	6	8	10	12	14	16
Global efficiency	100.00%	88.82%	84.18%	68.81%	55.16%	42.88%	36.28%	31.10%	27.12%
Parallelization strategy efficiency	100.00%	89.00%	84.80%	73.86%	59.41%	47.75%	40.49%	34.76%	30.64%
Load balancing	100.00%	89.03%	84.85%	73.95%	59.49%	47.82%	40.57%	34.83%	30.71%
In execution efficiency	100.00%	99.96%	99.94%	99.88%	99.86%	99.85%	99.80%	99.80%	99.79%
Scalability for computation tasks	100.00%	99.80%	99.27%	93.16%	92.84%	89.81%	89.61%	89.45%	88.49%
IPC scalability	100.00%	99.98%	99.93%	99.94%	99.93%	99.93%	99.91%	99.93%	99.93%
Instruction scalability	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%
Frequency scalability	100.00%	99.82%	99.34%	93.21%	92.90%	89.87%	89.69%	89.51%	88.55%

Table 2: Analysis done on Thu Nov 21 12:15:25 PM CET 2024, par1307

In the third table we see that the time per explicit task is high as expected. We can also notice that the synchronization overhead increases up to 227.11% which is not good.

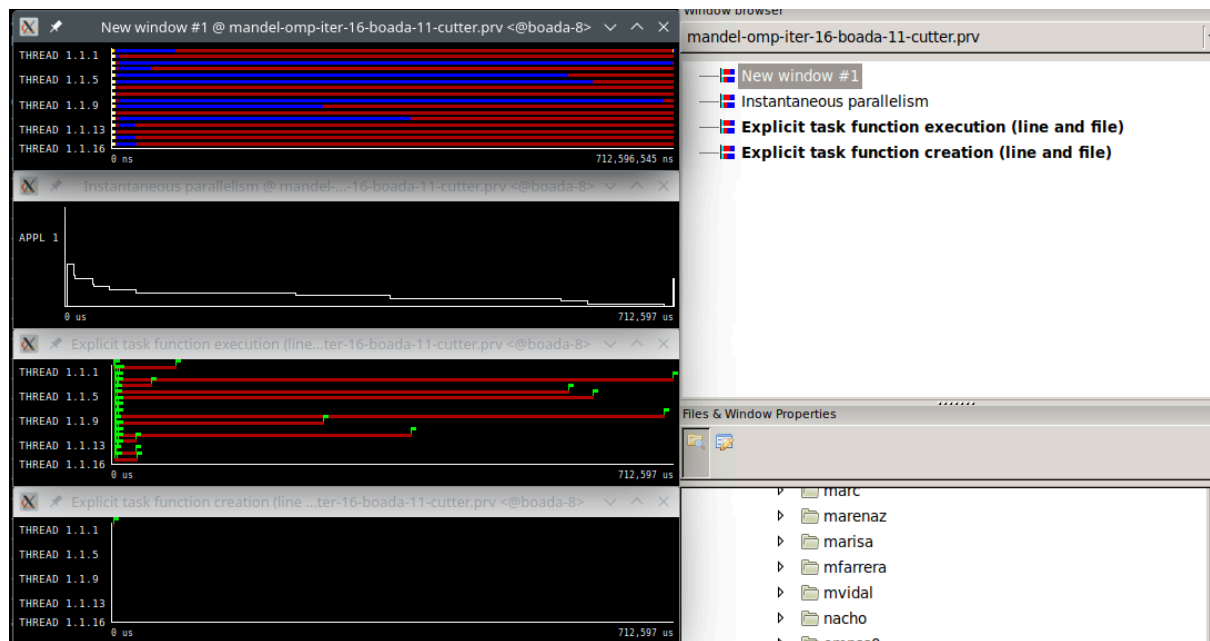
Statistics about explicit tasks in parallel fraction									
Number of Processors	1	2	4	6	8	10	12	14	16
Number of explicit tasks executed (total)	64.0	64.0	64.0	64.0	64.0	64.0	64.0	64.0	64.0
LB (number of explicit tasks executed)	1.0	0.94	0.64	0.43	0.47	0.38	0.31	0.27	0.24
LB (time executing explicit tasks)	1.0	0.89	0.85	0.74	0.59	0.48	0.41	0.35	0.31
Time per explicit task (average us)	48281.47	48368.7	48620.97	51792.18	51937.49	53632.85	53720.96	53739.44	54209.95
Overhead per explicit task (synch %)	0.0	12.33	17.87	35.28	68.18	109.4	146.93	187.93	227.11
Overhead per explicit task (sched %)	0.0	0.01	0.01	0.0	0.0	0.0	0.01	0.01	0.01
Number of taskwait/taskgroup (total)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Table 3: Analysis done on Thu Nov 21 12:15:25 PM CET 2024, par1307

Paraver Analysis

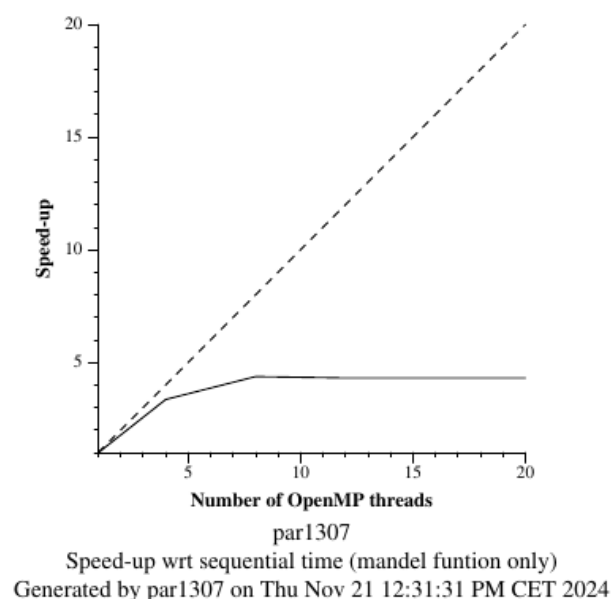
In the following picture we can see different paraver windows: execution trace, instantaneous parallelism and the explicit task function execution and creation.

Analysing the different images we can notice that, as mentioned before, parallelism is exploited at the beginning, but there is a point, when the processors finish their tasks the parallelism drops instantaneously.



Strong Scalability

This graphics shows the speed-up against the number of OpenMP threads, we notice that when using 1 to 4 threads it is close to the ideal speed-up, but as we keep increasing the threads, it starts getting constant, tending to a speed-up of 5.



	Number of threads					
Iterative: Tile	1	4	8	12	16	20
Elapsed time (ns)	3.089155	0.910804	0.698836	0.709182	0.709561	0.709295

Table 4: Elapsed time in (ns) when computing the iterative tile version.

Finer

Code

For the finer grain implementation we added some modifications to the tile approach based on the analysis done in the previous laboratory assignment:

We divided the `equal` variable in `equal1` and `equal2`, then in order to create the tasks, we added a `#pragma omp task shared(equal1)` and `#pragma omp task shared(equal2)` directives for each loop (vertical and horizontal borders) together with a `#pragma omp taskwait` to avoid data races. Finally, we will find some `#pragma omp task firstprivate(...)` directives which will create the finer grain tasks, one for each computation loop.

```

145 void mandel_tiled(int M[ROWS][COLS], double CminR, double CminI, double CmaxR, double CmaxI, double scale_real, double scale_imag, int maxiter)
146 {
147     int equal1 = 0, equal2 = 0;
148
149     for (int y = 0; y < ROWS; y += TILE)
150     for (int x = 0; x < COLS; x += TILE)
151     {
152         equal1 = 1;
153         equal2 = 1;
154         #pragma omp task firstprivate (x,y)
155         {
156             #pragma omp task shared(equal1)
157             // Set all matrix positions with the same value
158             // equal = 1;
159             for (int px = x; px < x + TILE; px++) {
160                 M[y][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y, scale_real, scale_imag, maxiter);
161                 M[y + TILE - 1][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y + TILE - 1, scale_real, scale_imag, maxiter);
162                 equal1 = equal1 & (M[y][x] == M[y][px]);
163                 equal1 = equal1 & (M[y][x] == M[y + TILE - 1][px]);
164             }
165             #pragma omp task shared(equal2)
166             for (int py = y; py < y + TILE; py++) {
167                 M[py][x] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x, py, scale_real, scale_imag, maxiter);
168                 M[py][x + TILE - 1] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x + TILE - 1, py, scale_real, scale_imag, maxiter);
169                 equal2 = equal2 & (M[y][x] == M[py][x]);
170                 equal2 = equal2 & (M[y][x] == M[py][x + TILE - 1]);
171             }
172
173             #pragma omp taskwait
174
175             #pragma omp task firstprivate (equal1, equal2, x, y)
176             {
177                 if (equal1 && equal2 && M[y][x] == maxiter) {
178                     if (output2histogram) {
179                         #pragma omp atomic
180                         histogram[M[y][x] - 1] += (TILE * TILE);
181                     }
182
183                     long color = (long)((M[y][x] - 1) * scale_color) + min_color;
184
185                     for (int py = y; py < y + TILE; py++) {
186                         #pragma omp task firstprivate (x, y, py)
187                         for (int px = x; px < x + TILE; px++) {
188                             M[py][px] = M[y][x];
189                             if (output2display) {
190                                 /* Scale color and display point */
191                                 if (setup_return == EXIT_SUCCESS) {
192                                     #pragma omp critical
193                                     {
194                                         XSetForeground(display, gc, color);
195                                         XDrawPoint(display, win, gc, px, py);
196                                     }
197                                 }
198                             }
199                         }
200                     }
201                 }
202             }
203         }
204     }
205
206     // Compute
207     for (int py = y; py < y + TILE; py++) {
208         #pragma omp task firstprivate (x, y, py)
209         for (int px = x; px < x + TILE; px++) {
210             M[py][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, py, scale_real, scale_imag, maxiter);
211             if (output2histogram) {
212                 #pragma omp atomic
213                 histogram[M[py][px] - 1]++;
214             }
215             if (output2display) {
216                 /* Scale color and display point */
217                 long color = (long)((M[py][px] - 1) * scale_color) + min_color;
218                 if (setup_return == EXIT_SUCCESS) {
219                     #pragma omp critical
220                     {
221                         XSetForeground(display, gc, color);
222                         XDrawPoint(display, win, gc, px, py);
223                     }
224                 }
225             }
226         }
227     }
228 }
229
230
231

```


Modelfactor Analysis

These are the results from running the modelfactor analysis:

In this first table we can notice that with the finer grain strategy, the elapsed time keeps decreasing when increasing the number of processors. Same behaviour can be seen in the speedup row.

Overview of whole program execution metrics									
Number of Processors	1	2	4	6	8	10	12	14	16
Elapsed time (sec)	3.11	1.57	0.81	0.57	0.43	0.36	0.31	0.27	0.24
Speedup	1.00	1.98	3.84	5.48	7.16	8.56	10.15	11.67	13.04
Efficiency	1.00	0.99	0.96	0.91	0.90	0.86	0.85	0.83	0.82

Table 1: Analysis done on Thu Nov 21 01:44:46 PM CET 2024, par1307

In the second table we see that the efficiency keeps quite constant which means that it has a great scalability. In this case, the load balancing drops a bit when increasing the number of processors but it stills keep above 95%.

Overview of the Efficiency metrics in parallel fraction, $\phi=99.99\%$									
Number of Processors	1	2	4	6	8	10	12	14	16
Global efficiency	99.92%	98.85%	96.05%	91.36%	89.53%	85.66%	84.62%	83.42%	81.57%
Parallelization strategy efficiency	99.92%	99.06%	98.22%	97.64%	96.18%	95.16%	94.31%	92.96%	91.35%
Load balancing	100.00%	99.91%	99.62%	99.41%	99.07%	98.35%	97.99%	96.25%	95.65%
In execution efficiency	99.92%	99.15%	98.59%	98.22%	97.08%	96.76%	96.24%	96.59%	95.50%
Scalability for computation tasks	100.00%	99.79%	97.79%	93.56%	93.08%	90.01%	89.73%	89.73%	89.30%
IPC scalability	100.00%	99.85%	99.87%	99.91%	99.80%	99.83%	99.83%	99.80%	99.81%
Instruction scalability	100.00%	100.10%	100.10%	100.10%	100.10%	100.10%	100.09%	100.10%	100.09%
Frequency scalability	100.00%	99.84%	97.83%	93.56%	93.18%	90.07%	89.80%	89.83%	89.38%

Table 2: Analysis done on Thu Nov 21 01:44:46 PM CET 2024, par1307

In the third table we see that the time per explicit task increases a bit together with the processors and the synchronization overhead just increases up to 4.84%.

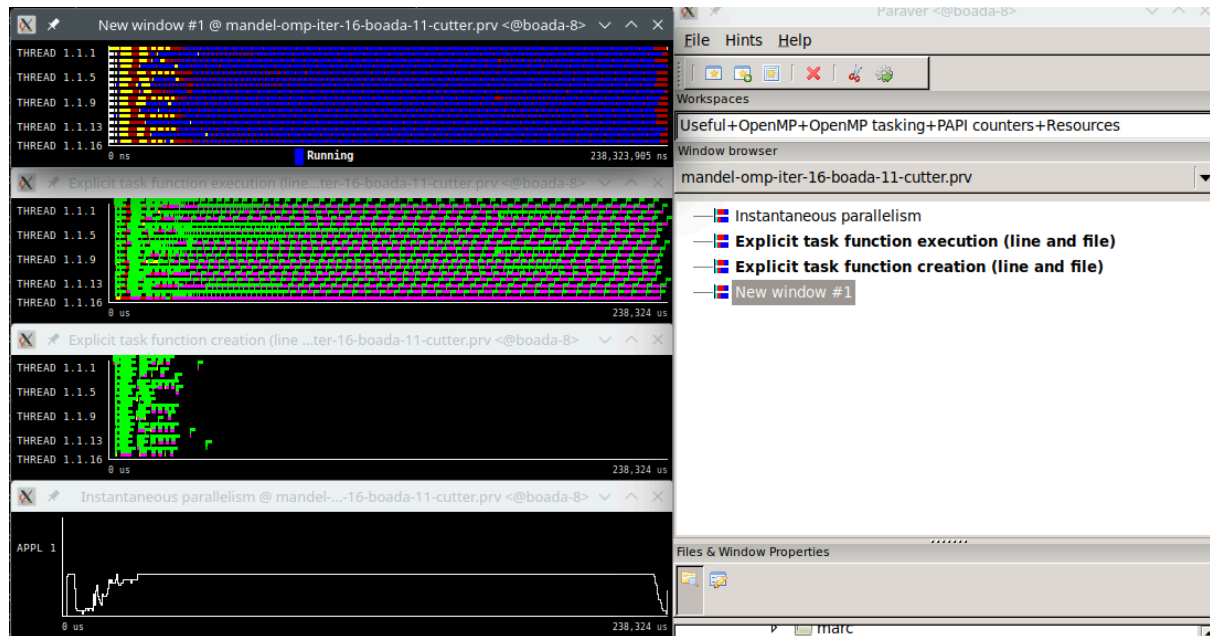
Statistics about explicit tasks in parallel fraction									
Number of Processors	1	2	4	6	8	10	12	14	16
Number of explicit tasks executed (total)	8448.0	8448.0	8448.0	8448.0	8448.0	8448.0	8448.0	8448.0	8448.0
LB (number of explicit tasks executed)	1.0	0.95	0.64	0.43	0.45	0.37	0.31	0.34	0.37
LB (time executing explicit tasks)	1.0	1.0	1.0	1.0	0.99	0.99	0.99	0.98	0.97
Time per explicit task (average us)	367.01	369.52	378.16	396.37	400.24	414.76	417.52	419.97	425.12
Overhead per explicit task (synch %)	0.0	0.59	1.09	1.29	2.17	2.84	3.31	4.01	4.84
Overhead per explicit task (sched %)	0.08	0.34	0.65	0.98	1.51	1.8	2.18	2.89	3.64
Number of taskwait/taskgroup (total)	64.0	64.0	64.0	64.0	64.0	64.0	64.0	64.0	64.0

Table 3: Analysis done on Thu Nov 21 01:44:46 PM CET 2024, par1307

Paraver Analysis

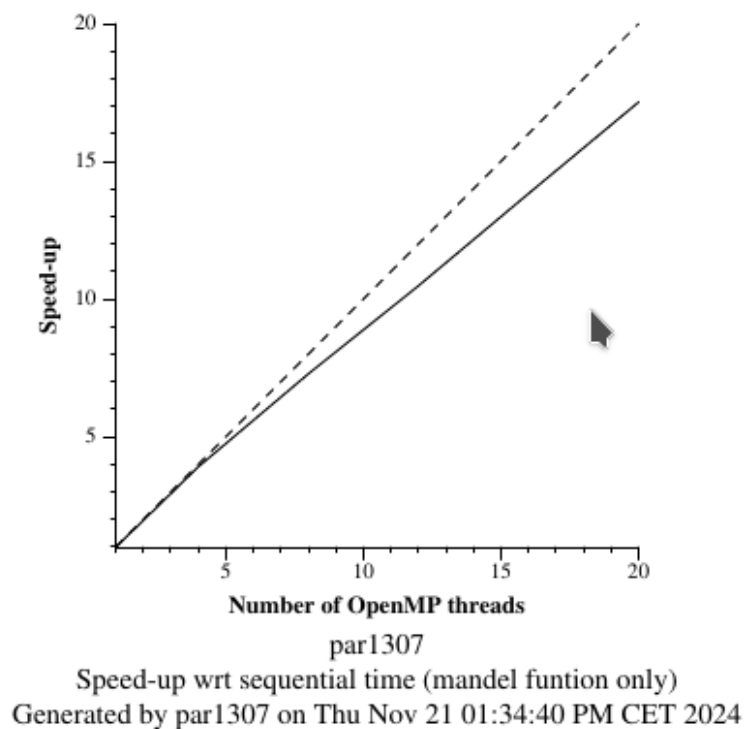
In the following picture we can see different paraver windows: execution trace, instantaneous parallelism and the explicit task function execution and creation.

Before we saw that parallelism was only exploited at the beginning, but now we see that parallelism is all over the execution.



Strong Scalability

This graphics shows the speed-up against the number of OpenMP threads, we notice that this strategy's behaviour is almost the ideal, meaning that it indeed has a strong scalability.



	Number of threads					
Iterative: Finer Grain	1	4	8	12	16	20
Elapsed time (ns)	3.094647	0.780898	0.418766	0.291014	0.220544	0.177836

Table 4: Elapsed time in (ns) when computing the iterative finer version.

Recursive task decomposition

Leaf

Code

```

145 void mandel_tiled_rec(int M[ROWS][COLS], int NRows, int NCols, int start_fil, int start_col, double CminR, double CminI, double CmaxR, double CmaxI, double scale_real, double scale_imag, int maxiter)
146 {
147     int non_first_call = (NRows != ROWS) ? (NCols != COLS) : 0;
148     int equal;
149     int x, y;
150
151     equal = non_first_call;
152     y = start_fil;
153     x = start_col;
154     for (int px = x; px < x + NCols; px++) {
155         M[y][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y, scale_real, scale_imag, maxiter);
156         M[y + NRows - 1][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, y + NRows - 1, scale_real, scale_imag, maxiter);
157         equal = equal & (M[y][x] == M[y][px]);
158         equal = equal & (M[y][x] == M[y + NRows - 1][px]);
159     }
160     for (int py = y; py < y + NRows; py++) {
161         M[py][x] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x, py, scale_real, scale_imag, maxiter);
162         M[py][x + NCols - 1] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, x + NCols - 1, py, scale_real, scale_imag, maxiter);
163         equal = equal & (M[y][x] == M[py][x]);
164         equal = equal & (M[y][x] == M[py][x + NCols - 1]);
165     }
166
167     //check if we can avoid computation of tile
168     if (equal) {
169         // Set all matrix positions with the same value
170         #pragma omp task
171         {
172             long color = (long)((M[y][x] - 1) * scale_color) + min_color;
173             if (output2histogram) {
174                 #pragma omp atomic
175                 histogram[M[y][x] - 1] += (NRows * NCols);
176             }
177             for (int py = y; py < y + NRows; py++) {
178                 for (int px = x; px < x + NCols; px++) {
179                     M[py][px] = M[y][x];
180                     if (output2display) {
181                         if (setup_return == EXIT_SUCCESS) {
182                             #pragma omp critical
183                             {
184                                 XSetForeground(display, gc, color);
185                                 XDrawPoint(display, win, gc, px, py);
186                             }
187                         }
188                     }
189                 }
190             }
191         }
192     } else {
193         if (NCols <= TILE) {
194             // Compute
195             #pragma omp task firstprivate(y,x)
196             for (int py = y; py < y + NRows; py++) {
197                 for (int px = x; px < x + NCols; px++) {
198                     M[py][px] = pixel_dwell(COLS, ROWS, CminR, CminI, CmaxR, CmaxI, px, py, scale_real, scale_imag, maxiter);
199                     if (output2histogram) {
200                         #pragma omp atomic
201                         histogram[M[py][px] - 1]++;
202                     }
203                 }
204             }
205             if (output2display) {
206                 /* Scale color and display point */
207                 long color = (long)((M[py][px] - 1) * scale_color) + min_color;
208                 if (setup_return == EXIT_SUCCESS) {
209                     #pragma omp critical
210                     {
211                         XSetForeground(display, gc, color);
212                         XDrawPoint(display, win, gc, px, py);
213                     }
214                 }
215             }
216         } else {
217             if (NRows > TILE) {
218                 mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
219                 mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
220                 mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil + NRows / 2, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
221                 mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil + NRows / 2, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
222             } else {
223                 mandel_tiled_rec(M, NRows, NCols / 2, start_fil, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
224                 mandel_tiled_rec(M, NRows, NCols / 2, start_fil, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
225             }
226         }
227     }
228 }

```

We added the leaf calls when reaching the last recursivity level (base case).

Modelfactor Analysis

Overview of whole program execution metrics									
Number of Processors	1	2	4	6	8	10	12	14	16
Elapsed time (sec)	1.62	1.24	1.32	1.33	1.37	1.37	1.37	1.37	1.37
Speedup	1.00	1.30	1.22	1.22	1.18	1.18	1.18	1.18	1.18
Efficiency	1.00	0.65	0.31	0.20	0.15	0.12	0.10	0.08	0.07

Table 1: Analysis done on Thu Nov 28 12:42:51 PM CET 2024, par1307

Lab 4: Parallel Task Decomposition Implementation and Analysis

Overview of the Efficiency metrics in parallel fraction, $\phi=99.98\%$									
Number of Processors	1	2	4	6	8	10	12	14	16
Global efficiency	99.94%	64.97%	30.56%	20.31%	14.73%	11.76%	9.82%	8.40%	7.35%
Parallelization strategy efficiency	99.94%	65.20%	32.39%	21.79%	16.27%	13.10%	10.93%	9.38%	8.25%
Load balancing	100.00%	65.43%	32.51%	21.87%	16.34%	13.15%	10.97%	9.42%	8.28%
In execution efficiency	99.94%	99.66%	99.63%	99.63%	99.61%	99.59%	99.60%	99.60%	99.61%
Scalability for computation tasks	100.00%	99.64%	94.34%	93.21%	90.51%	89.81%	89.79%	89.53%	89.20%
IPC scalability	100.00%	99.66%	99.61%	99.53%	99.53%	99.46%	99.55%	99.52%	99.50%
Instruction scalability	100.00%	100.06%	100.07%	100.06%	100.06%	100.06%	100.06%	100.06%	100.06%
Frequency scalability	100.00%	99.91%	94.65%	93.59%	90.88%	90.24%	90.15%	89.91%	89.59%

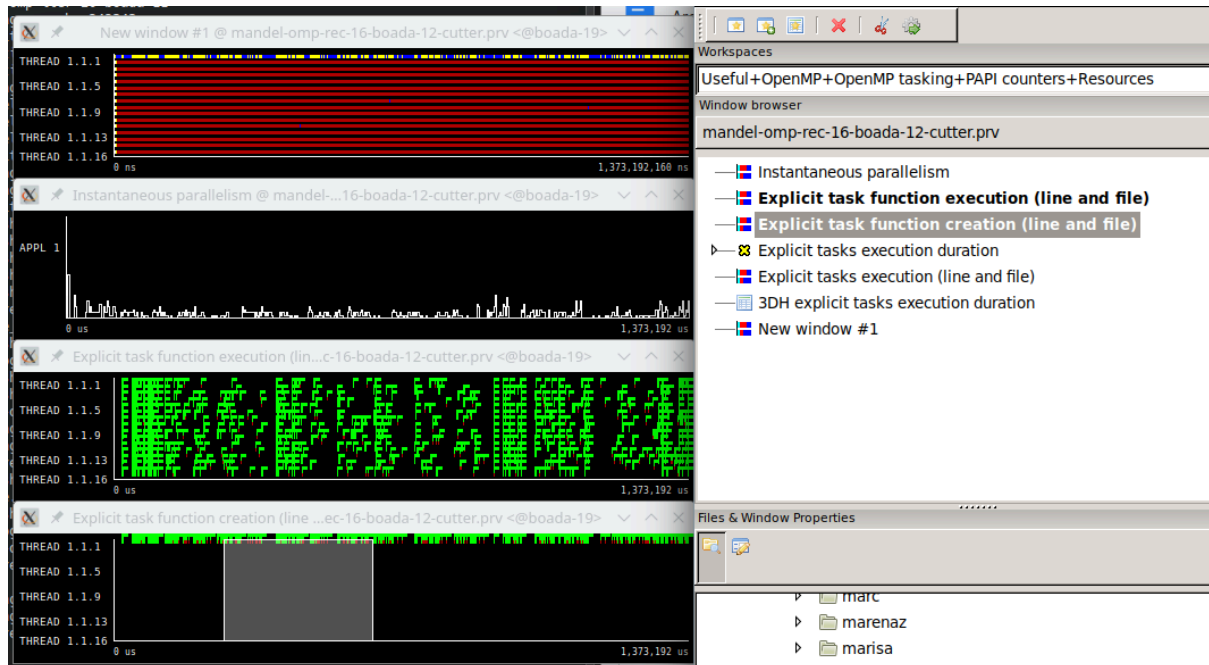
Table 2: Analysis done on Thu Nov 28 12:42:51 PM CET 2024, par1307

Statistics about explicit tasks in parallel fraction									
Number of Processors	1	2	4	6	8	10	12	14	16
Number of explicit tasks executed (total)	3013.0	3013.0	3013.0	3013.0	3013.0	3013.0	3013.0	3013.0	3013.0
LB (number of explicit tasks executed)	1.0	0.56	0.65	0.89	0.8	0.87	0.78	0.79	0.81
LB (time executing explicit tasks)	1.0	0.5	0.67	0.9	0.81	0.7	0.64	0.6	0.76
Time per explicit task (average us)	125.93	127.17	130.71	135.99	137.67	140.69	140.55	140.43	140.56
Overhead per explicit task (synch %)	0.0	224.8	905.69	1516.49	2210.39	2810.39	3456.56	4113.86	4752.27
Overhead per explicit task (sched %)	0.23	0.82	1.05	1.02	0.97	0.94	0.98	0.95	0.91
Number of taskwait/taskgroup (total)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Table 3: Analysis done on Thu Nov 28 12:42:51 PM CET 2024, par1307


Thanks to the Modelfactor tables, we can observe a point (8 processors) where the speedup reaches its peak. Moreover, efficiency drops sharply beyond 4 processors, which could be attributed to poor load balancing management. This is confirmed in the third table, where we can see there is indeed significant load imbalance.

Paraver Analysis



In this execution of wxparaver over this particular parallel implementation, we can see all the tasks created at leaf level. The explicit task execution graph shows a huge Load Unbalance, as there are a lot of spaces where threads are stopped waiting for other threads to finish. This is not efficient.

Strong Scalability



Speed-up wrt sequential time (mandel function only)

Generated by par1307 on Thu Nov 28 12:49:07 PM CET 2024

1	1.611466
4	1.242143
8	1.323873
12	1.368214
16	1.369712
20	1.370270

Using the strong scalability analysis, however, we can observe severe efficiency problems, which might be caused by a poor parallelisation strategy, big Load Unbalance or Memory Contention problems.

Tree

Code

```

214 } else {
215     if (NRows > TILE) {
216         #pragma omp task
217         mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
218         #pragma omp task
219         mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
220         #pragma omp task
221         mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil + NRows / 2, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
222         #pragma omp task
223         mandel_tiled_rec(M, NRows / 2, NCols / 2, start_fil + NRows / 2, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
224     } else {
225         #pragma omp task
226         mandel_tiled_rec(M, NRows, NCols / 2, start_fil, start_col, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
227         #pragma omp task
228         mandel_tiled_rec(M, NRows, NCols / 2, start_fil, start_col + NCols / 2, CminR, CminI, CmaxR, CmaxI, scale_real, scale_imag, maxiter);
229     }
230 }
231 }
232 }
233 }

```

We added the recursive call tasks, which create the Tree shaped core of this parallelisation strategy.

Modelfactor Analysis

Overview of whole program execution metrics									
Number of Processors	1	2	4	6	8	10	12	14	16
Elapsed time (sec)	1.62	0.82	0.44	0.32	0.25	0.22	0.19	0.17	0.16
Speedup	1.00	1.97	3.67	5.10	6.39	7.42	8.57	9.48	10.40
Efficiency	1.00	0.99	0.92	0.85	0.80	0.74	0.71	0.68	0.65

Table 1: Analysis done on Thu Nov 28 12:58:20 PM CET 2024, par1307

Overview of the Efficiency metrics in parallel fraction, $\phi=99.98\%$									
Number of Processors	1	2	4	6	8	10	12	14	16
Global efficiency	99.88%	98.54%	91.61%	85.03%	79.88%	74.18%	71.39%	67.78%	65.05%
Parallelization strategy efficiency	99.88%	98.53%	94.63%	90.97%	87.00%	82.36%	79.40%	75.73%	72.95%
Load balancing	100.00%	99.30%	96.46%	96.64%	95.00%	93.02%	90.15%	89.02%	84.76%
In execution efficiency	99.88%	99.22%	98.10%	94.13%	91.58%	88.54%	88.07%	85.07%	86.06%
Scalability for computation tasks	100.00%	100.00%	96.81%	93.47%	91.81%	90.07%	89.92%	89.51%	89.18%
IPC scalability	100.00%	99.86%	99.83%	99.76%	99.68%	99.60%	99.60%	99.50%	99.51%
Instruction scalability	100.00%	100.16%	100.16%	100.16%	100.16%	100.16%	100.15%	100.15%	100.15%
Frequency scalability	100.00%	99.99%	96.82%	93.55%	91.96%	90.29%	90.14%	89.82%	89.48%

Table 2: Analysis done on Thu Nov 28 12:58:20 PM CET 2024, par1307

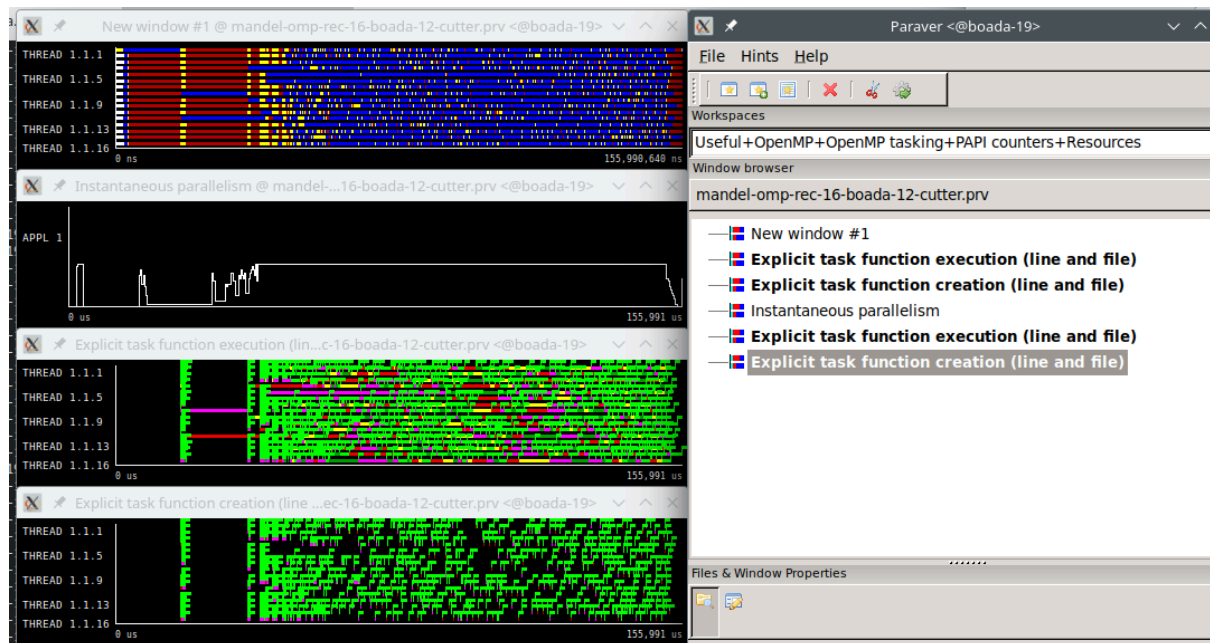
Statistics about explicit tasks in parallel fraction									
Number of Processors	1	2	4	6	8	10	12	14	16
Number of explicit tasks executed (total)	7029.0	7029.0	7029.0	7029.0	7029.0	7029.0	7029.0	7029.0	7029.0
LB (number of explicit tasks executed)	1.0	0.78	0.6	0.61	0.68	0.45	0.55	0.49	0.64
LB (time executing explicit tasks)	1.0	1.0	0.98	0.96	0.94	0.92	0.89	0.89	0.85
Time per explicit task (average us)	228.07	228.64	236.32	244.71	249.29	254.62	255.34	256.98	257.74
Overhead per explicit task (synch %)	0.0	1.3	5.37	9.52	14.26	20.24	24.21	29.56	33.86
Overhead per explicit task (sched %)	0.12	0.16	0.23	0.27	0.44	0.74	1.02	1.46	1.76
Number of taskwait/taskgroup (total)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Table 3: Analysis done on Thu Nov 28 12:58:20 PM CET 2024, par1307

In these tables we can see a much better speedup compared to the previous strategy, as it doesn't get stuck at all when using an increasingly big number of processors. Furthermore, the efficiency of the parallelization strategy remains over 70% even at a high number of processors.

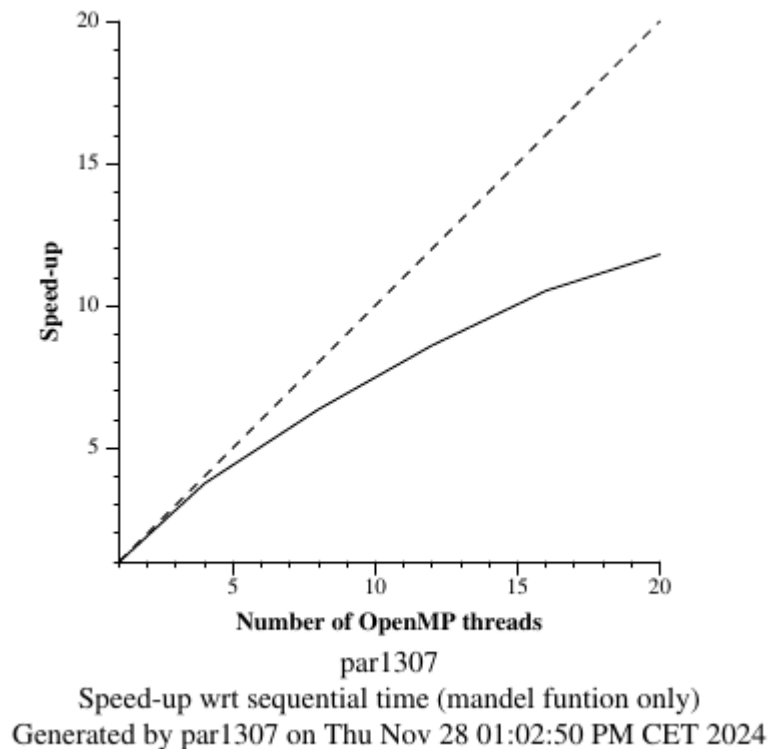
Furthermore and regarding the efficiency of this parallelisation strategy, we can see that it drops to ~70% when using 20 threads, but that's not too bad.

Paraver Analysis



In this execution of wxparaver over this particular parallel implementation, we can see all the tasks created at each tree level. All the graphs look reasonably balanced, therefore there is not an apparent big Load Unbalance or task creation overhead problem causing a bottleneck that can be seen by the naked eye on these graphs.

Strong Scalability



Lab 4: Parallel Task Decomposition Implementation and Analysis

1	1.612781
4	0.427873
8	0.252638
12	0.186513
16	0.152506
20	0.136090

In this graph we can see how the speedup has more of a logarithmic shape, which is expected, but the efficiency and overall health of the speedup evolution is much better than that of the previous strategy.

Summary of the elapsed execution time

	Number of threads					
Version	1	4	8	12	16	20
Iterative: Tile	3.089155	0.910804	0.698836	0.709182	0.709561	0.709295
Iterative: Finer Grain	3.094647	0.780898	0.418766	0.291014	0.220544	0.177836
Recursive: Leaf	1.611466	1.242143	1.323873	1.368214	1.369712	1.370270
Recursive: Tree	1.612781	0.427873	0.252638	0.186513	0.152506	0.136090

Summary of the elapsed execution times for each of the versions, obtained from the output files after the execution of `submit-strong-omp.sh` script